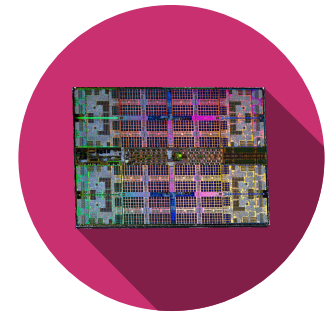


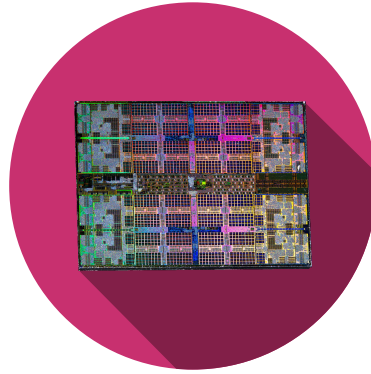


# Computer Architecture Performance Per

Information and Communication Systems program


Silvan Zahno [silvan.zahno@hevs.ch](mailto:silvan.zahno@hevs.ch)





Benchmarks

# Benchmarks




Judges the relative performance of a CPU

## **Key terms**

- MIPS
- MFLOPS
- MHz
- FPS
- Render Time
- Dropped Frames

# Benchmarks

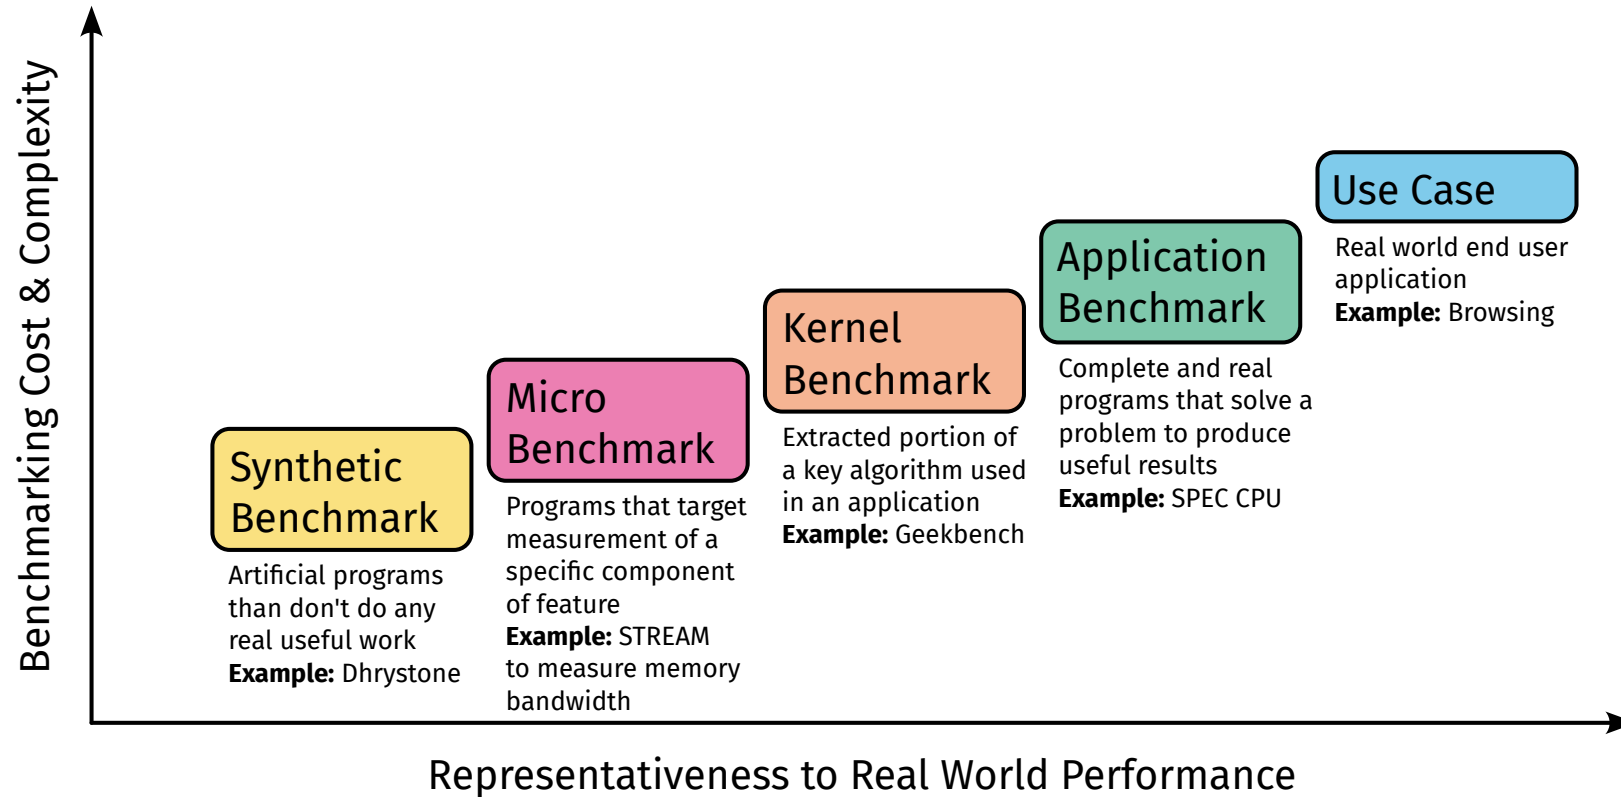



Judges the relative performance of a CPU

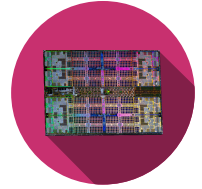
## Key terms

- MIPS – Million Instruction per Second
- MFLOPS – Million Floating Point Operation per Second
- MHz - Frequency
- FPS – Frame per second
- Render Time – Time to render 2D/3D scene
- Dropped Frames – number of frames lost in streaming


# Spectrum of Benchmarks




# SPEC Benchmark SPEC CPU 2017



- **Standard Performance Evaluation Corporation**
- Measures integer and floating point operation performance
- Contains 10 integer and 13 floating point operations
- Compute intensive, concentrates on CPU and memory




# Linpack




## Performance per Watt

- Performances of double precision vector and matrix operations
- Often used by HPC community to measure FLOPS of a processor



# Aspects of Computer Performance



## Time


- Response time / latency / execution time
  - Time between start and completion of event / task / program (n seconds)

## Other Aspects

- Availability, scalability, performance per watt




# Benchmark Suites



Benchmark	Organization	Description	Comment
Dhrystone	Open Source	Synthetic benchmark measure compute-intensive integer performance	No longer representative of modern workloads
CoreMark	EEMBC	Popular benchmark to measure compute-intensive integer performance	Result published through EEMBS, continuation of Dhrystone
SPEC CPU	SPEC OSG CPU	Application benchmarks that comprise of open-source real world applications.	Most popular for measuring CPU performance
STREAM	Open Source	Suite of microbenchmarks to measure sustained memory bandwidth	Measure of memory latency and bandwidth
Geekbench 5	Primate labs	Kernel benchmarks that measure CPU Integer, Floating Point and Memory performance	Used in industry, but with some limitation on the representation of real-world scenarios
PassMark	PassMark Software	Kernel benchmarks that measure CPU Integer, Floating Point and Memory performance Per	Used in industry, but with some limitation on the representation of real-world scenarios

# SPEC Benchmarks



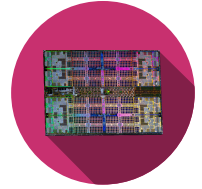
SPECrate@2017 Integer	SPECspeed@2017 Integer	Language [1]	KLOC [2]	Application Area
<a href="#">500.perlbench_r</a>	<a href="#">600.perlbench_s</a>	C	362	Perl interpreter
<a href="#">502.gcc_r</a>	<a href="#">602.gcc_s</a>	C	1,304	GNU C compiler
<a href="#">505.mcf_r</a>	<a href="#">605.mcf_s</a>	C	3	Route planning
<a href="#">520.omnetpp_r</a>	<a href="#">620.omnetpp_s</a>	C++	134	Discrete Event simulation - computer network
<a href="#">523.xalancbmk_r</a>	<a href="#">623.xalancbmk_s</a>	C++	520	XML to HTML conversion via XSLT
<a href="#">525.x264_r</a>	<a href="#">625.x264_s</a>	C	96	Video compression
<a href="#">531.deepsjeng_r</a>	<a href="#">631.deepsjeng_s</a>	C++	10	Artificial Intelligence: alpha-beta tree search (Chess)
<a href="#">541.leela_r</a>	<a href="#">641.leela_s</a>	C++	21	Artificial Intelligence: Monte Carlo tree search (Go)
<a href="#">548.exchange2_r</a>	<a href="#">648.exchange2_s</a>	Fortran	1	Artificial Intelligence: recursive solution generator (Sudoku)
<a href="#">557.xz_r</a>	<a href="#">657.xz_s</a>	C	33	General data compression

SPECrate@2017 Floating Point	SPECspeed@2017 Floating Point	Language [1]	KLOC [2]	Application Area
<a href="#">503.bwaves_r</a>	<a href="#">603.bwaves_s</a>	Fortran	1	Explosion modeling
<a href="#">507.cactuBSSN_r</a>	<a href="#">607.cactuBSSN_s</a>	C++, C, Fortran	257	Physics: relativity
<a href="#">508.namd_r</a>		C++	8	Molecular dynamics
<a href="#">510.parest_r</a>		C++	427	Biomedical imaging: optical tomography with finite elements
<a href="#">511.povray_r</a>		C++, C	170	Ray tracing
<a href="#">519.lbm_r</a>	<a href="#">619.lbm_s</a>	C	1	Fluid dynamics
<a href="#">521.wrf_r</a>	<a href="#">621.wrf_s</a>	Fortran, C	991	Weather forecasting
<a href="#">526.blender_r</a>		C++, C	1,577	3D rendering and animation
<a href="#">527.cam4_r</a>	<a href="#">627.cam4_s</a>	Fortran, C	407	Atmosphere modeling
	<a href="#">628.pop2_s</a>	Fortran, C	338	Wide-scale ocean modeling (climate level)
<a href="#">538.imagick_r</a>	<a href="#">638.imagick_s</a>	C	259	Image manipulation
<a href="#">544.nab_r</a>	<a href="#">644.nab_s</a>	C	24	Molecular dynamics
<a href="#">549.fotonik3d_r</a>	<a href="#">649.fotonik3d_s</a>	Fortran	14	Computational Electromagnetics
<a href="#">554.roms_r</a>	<a href="#">654.roms_s</a>	Fortran	210	Regional ocean modeling

[1] For multi-language benchmarks, the first one listed determines library and link options ([details](#))

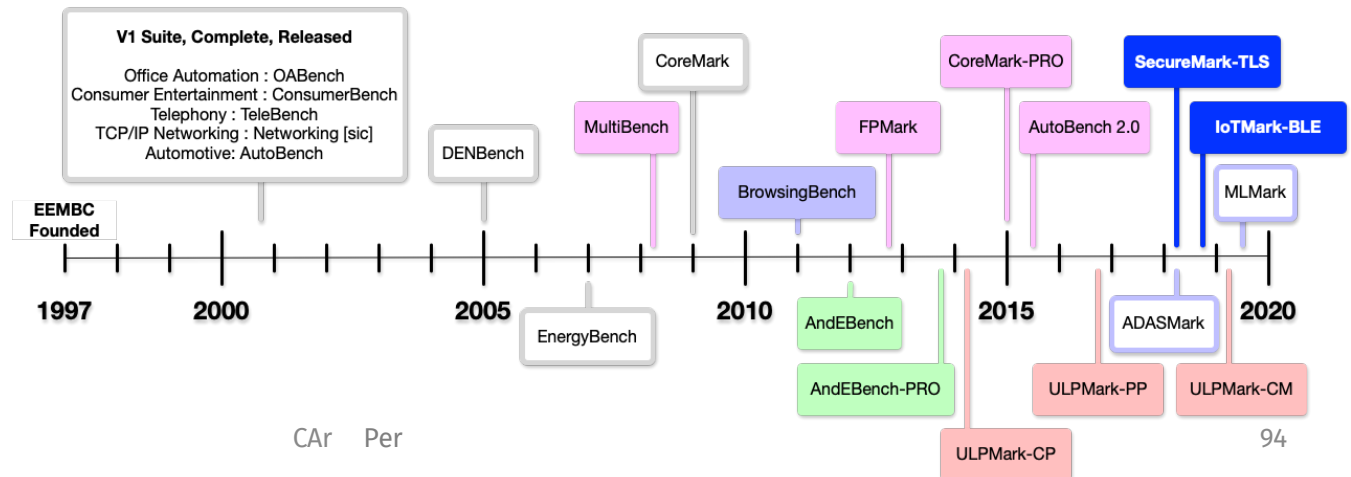
[2] KLOC = line count (including comments/whitespace) for source files used in a build / 1000

# EEMBC Benchmarks




Benchmarks can be grouped into the following categories:

- [Ultra-Low Power and Internet of Things](#)
- [Heterogeneous Compute](#)
- [Single-core Processor Performance](#)
- [Multi-core Processor Performance](#)
- [Phone and Tablet](#)



# How to Measure Execution time?

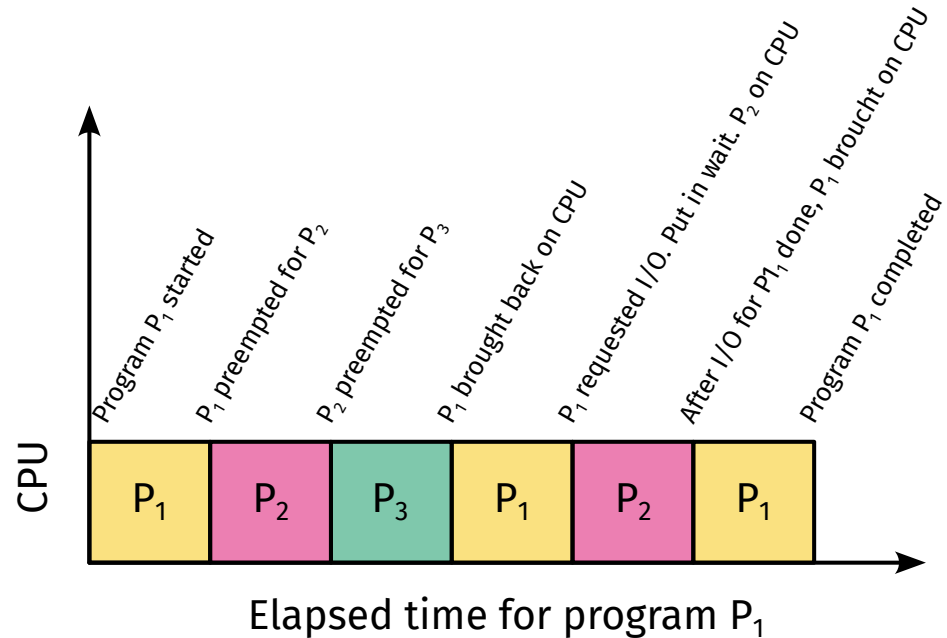


## Wall-clock time

- Elapsed time
  - Disk Access, I/O, OS overhead, ...


Multiprogrammed CPU works on other process when current process stalled for I/O

CPU-time = time CPU is computing



# Benchmark

## Exercise



Which of the following statements are correct?

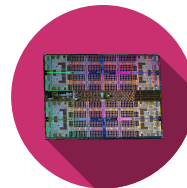
1. The wall clock time is the total elapsed time , including I/O, Operating system overhead etc.?
2. Multi-threading improves the throughput of a process
3. The CPU time does not include the I/O time
4. Multi-threading improves the execution time of a process

# Benchmark


## Quick and Dirty

For quickly benchmark your application or functions you can use hyperfine. A rust base CLI Program

<https://github.com/sharkdp/hyperfine>




```
► hyperfine --warmup 3 'fd -e jpg -uu' 'find -iname "*.jpg"'
```



Performance

# Summarizing Performance




- With 1 program it is clear which computer is faster
- With  $>1$  programs, it depends

How do you aggregate their performance?



# Summarizing Performance




- With 1 program it is clear which computer is faster
- With >1 programs, it depends

How do you aggregate their performance?

$$\textit{ArithmeticMean} = AM = \frac{1}{n} \sum_{i=1}^n \textit{Time}_i$$

# Summarizing Performance

## Exercise




	CPU A	CPU B	CPU C
Program $P_1$ (sec)	1	10	20
Program $P_2$ (sec)	1000	100	20

1. Which CPU is the fastest for  $P_1$  ?
2. Which CPU is the fastest for  $P_2$  ?
3. Which CPU is the fastest?

# Summarizing Performance


## Exercise



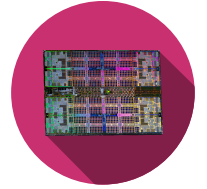
	CPU A	CPU B	CPU C
Program $P_1$ (sec)	1	10	20
Program $P_2$ (sec)	1000	100	20
Total (sec)			
Arithmetic Mean			

1. Which CPU is the fastest for  $P_1$  ?
2. Which CPU is the fastest for  $P_2$  ?
3. Which CPU is the fastest ?

# Summarizing Performance



- What if the program not run equally often?
- Two approaches:
  - Weighted execution time
  - Normalize execution times to a reference machine and take the average



# Summarizing Performance


## Weighted Execution Time

- Program  $P_i$  has a weight  $Weight_i$ 
  - Indicates execution frequency
  - weights add up to 1
- Program  $P_i$  takes time  $Time_i$
- Weighted arithmetic mean:

$$WeightedArithmeticMean = WAM = \sum_{i=1}^n Weight_i * Time_i$$

# Summarizing Performance

## Weighted Execution Time - Example




	CPU A	CPU B	CPU C
Program $P_1$ (sec)	1	10	20
Program $P_2$ (sec)	1000	100	20
Arithmetic mean	500.5	55	20

1. Which CPU is the fastest with  $w_1=0.8$  and  $w_2 = 0.2$
2. Which CPU is the fastest with  $w_1=0.9$  and  $w_2 = 0.1$

# Summarizing Performance

## Weighted Execution Time - Example




	CPU A	CPU B	CPU C
Program $P_1$ (sec)	1	10	20
Program $P_2$ (sec)	1000	100	20
Arithmetic mean	500.5	55	20
Weighted mean ( $w_1=0.8, w_2=0.2$ )			
Weighted mean ( $w_1=0.9, w_2=0.1$ )			

1. Which CPU is the fastest with  $w_1=0.8$  and  $w_2 = 0.2$
2. Which CPU is the fastest with  $w_1=0.9$  and  $w_2 = 0.1$

# Summarizing Performance

## Weighted Execution Time - Exercise



We want to buy a new computer. It will mostly run programs  $P_1$  and  $P_2$ .


1. When is CPU A the best buy?
2. When is CPU B the best buy?
3. When is CPU C the best buy?

	CPU A	CPU B	CPU C
Program $P_1$ (sec)	1	10	100
Program $P_2$ (sec)	100	10	1



# Summarizing Performance

## Normalized Execution Time




- Normalize the execution time to a reference computer

$$ExecutionTimeRatio = \frac{ExecutionTime_{reference}}{ExecutionTime_{new}}$$

- Used by SPEC Benchmarks, called SPECRatio

# Summarizing Performance




- Normalized Execution Time - Example

	CPU A	CPU B	CPU C
Program P <sub>1</sub> (sec)	1	10	20
Program P <sub>2</sub> (sec)	1000	100	20

	Normalized to A			Normalized to B		
	CPU A	CPU B	CPU C	CPU A	CPU B	CPU C
ETR P <sub>1</sub>						
ETR P <sub>2</sub>						

# Summarizing Performance

## Normalized Execution Time - Example



- Do NOT use arithmetic mean to average normalized execution times!!


	CPU A	CPU B	CPU C
Program P <sub>1</sub> (sec)	1	10	20
Program P <sub>2</sub> (sec)	1000	100	20

	Normalized to A			Normalized to B		
	CPU A	CPU B	CPU C	CPU A	CPU B	CPU C
ETR P <sub>1</sub>	1	0.1	0.05	10	1	0.5
ETR P <sub>2</sub>	1	10	50	0.1	1	5.0
Arithmetic mean	1	5.05	25.025	5.05	1	2.75

# Summarizing Performance

## Geometric Mean

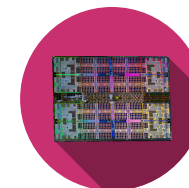
- Independent of reference
- No physical meaning



$$\textit{Geometric Mean} = GM = \sqrt[n]{\prod_{i=1}^n \textit{ExecutionTimeRatio}_i}$$

# Summarizing Performance

## Geometric Mean - Example




$$GM = \sqrt[n]{\prod_{i=1}^n ExecutionTimeRatio_i}$$

	CPU A	CPU B	CPU C
Program P <sub>1</sub> (sec)	1	10	20
Program P <sub>2</sub> (sec)	1000	100	20

	Normalized to A			Normalized to B		
	CPU A	CPU B	CPU C	CPU A	CPU B	CPU C
ETR P <sub>1</sub>	1	0.1	0.05	10	1	0.5
ETR P <sub>2</sub>	1	10	50	0.1	1	5.0
Geometric mean						

# CPU Performance Equation

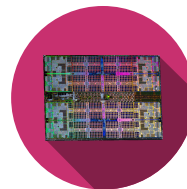


$$T = IC * CPI * CT = \frac{IC * CPI}{f}$$

- $T = \textit{Execution time}$
- $IC = N_{instr} = \# \text{ instructions executed (Instruction Count)}$
- $CPI = \textbf{C}ycles \textbf{P}er \textbf{I}nstruction$
- $CT = t_{cycle} = \textbf{C}ycle \textbf{T}ime = \text{duration of clock cycle}$
- $f = \text{clock frequency} = \frac{1}{t_{cycle}}$

# CPU Performance Equation

## Improving Performance




$$T = IC * CPI * CT = \frac{IC * CPI}{f}$$

- Factors are interdependent => Trade-offs
  - Cost
  - Power
  - Performance

Reduce	How	Side effect
$IC / N_{instr}$	Increase capability / complexity of instructions	$CPI / t_{cycle}$ increases
$CPI$	Simple instructions required fewer cycles	$IC / N_{instr}$ increases
$CT / t_{cycle}$	Less work per cycle	$CPI$ increases

# CPU Performance Equation

## Calculating Execution Time

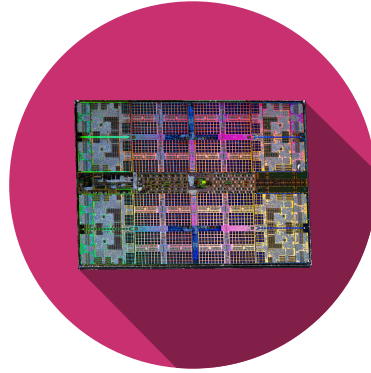


A program executes 5 Million instruction with a given instruction mix:

Instruction	Frequency	$CPI_{instr}$
ALU	50%	3
Load	20%	5
Store	10%	4
Branch	20%	3



- The CPU has a frequency of 2GHz
1. What is the execution time?








Amdahls Law





# Amdahl's Law







# Amdahl's Law



# Amdahl's Law








# Amdahl's Law








If 50% of the execution time is sequential, the maximum speedup is 2, no matter how many cores you use.  
- Gene Amdahl -

# Amdahl's Law



$$Time = s + \frac{f}{p}$$

# Amdahl's Law



$$S = \frac{1}{s + \frac{f}{p}} = \frac{1}{1 - f + \frac{f}{p}}$$

$S$  = Speedup

$s = 1 - f$  = serial fraction

$f = 1 - s$  = parallel fraction

$p$  = # processor cores

# Amdahl's Law

## Speedup for Different Parallel Fractions


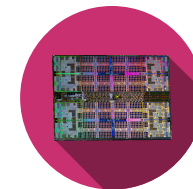




Image Credit: Wikipedia



# Amdahl's Law

## Generalization of Amdahl's Law

- Suppose we can improve fraction  $f$  by factor  $F$
- What will be overall improvement?



$$S = \frac{1}{s + \frac{f}{F}} = \frac{1}{1 - f + \frac{f}{F}}$$

$S$  = Speedup


$s = 1 - f$  = fraction we cannot improve

$f = 1 - s$  = fraction we can improve

$F$  = improvement factor

# Amdahl's Law


## Example



*FP instruction improved to run 2x as fast, but only 10% of all executed instructions are FP.*

- *What will be the overall speedup?*

# References



- [1]  
“11th Gen H Series: Benchmarks Across the Stack,” *Intel*. <https://www.intel.com/content/www/us/en/support/articles/000059752/processors.html> (accessed Jul. 21, 2022).
- [2]  
“Amdahl’s law,” *Wikipedia*. Jun. 29, 2022. Accessed: Jul. 18, 2022. [Online]. Available: [https://en.wikipedia.org/w/index.php?title=Amdahl%27s\\_law&oldid=1095605183](https://en.wikipedia.org/w/index.php?title=Amdahl%27s_law&oldid=1095605183)
- [3]  
ARM Ltd., “An Introduction to CPU Performance Benchmarks and How This Applies to the Home Market.” 2021.
- [4]  
S. H. Fuller and L. I. Millett, “Computing Performance: Game Over or Next Level?,” *Computer*, vol. 44, no. 1, pp. 31–38, Jan. 2011, doi: [10.1109/MC.2011.15](https://doi.org/10.1109/MC.2011.15).
- [5]  
“EEMBC.” <https://www.eembc.org/products/> (accessed Feb. 23, 2023).
- [6]  
C. benchmarks are crucial in determining whether your processor can run the games and applications you like H. what you need to get started.1, “How to Read and Understand CPU Benchmarks,” *Intel*. <https://www.intel.com/content/www/us/en/gaming/resources/read-cpu-benchmarks.html> (accessed Jul. 21, 2022).
- [7]  
“linpack performance per wat at DuckDuckGo.”  
<https://duckduckgo.com/?q=linpack+performance+per+wat&t=osx&iar=images&iax=images&ia=images&iai=http%3A%2F%2Fwww.intel.com%2Fcontent%2Fdam%2Fwww%2Fpublic%2Fus%2Fen%2Fimages%2Fcharts%2Fchart-id-642.png> (accessed Jul. 21, 2022).
- [8]  
“Overview - CPU 2017.” <https://www.spec.org/cpu2017/Docs/overview.html#benchmarks> (accessed Feb. 23, 2023).
- [9]  
“SPEC CPU2017 Results.” <https://www.spec.org/cpu2017/results/> (accessed Jul. 21, 2022).
- [10]  
*TU Berlin: Course AES*, (Oct. 22, 2018). Accessed: Jul. 18, 2022. [Online Video]. Available: <https://www.youtube.com/channel/UCPSsA8oxISBjidJsSPdpjsQ>


WHY ARE THERE MIRRORS ABOVE BEDS

WHY DO I SAY UH  
WHY IS SEA SALT BETTER

WHY ARE THERE TREES IN THE MIDDLE OF FIELDS  
WHY IS THERE NOT A POKEMON MMO  
WHY IS THERE LAUGHING IN TV SHOWS  
WHY ARE THERE DOORS ON THE FREEWAY  
WHY ARE THERE SO MANY SVCHOST:EXE RUNNING  
WHY AREN'T ANY COUNTRIES IN ANTARCTICA  
WHY ARE THERE SCARY SOUNDS IN MINECRAFT  
WHY IS THERE KICKING IN MY STOMACH  
WHY ARE THERE TWO SLASHES AFTER HTTP  
WHY ARE THERE CELEBRITIES  
WHY DO SNAKES EXIST  
WHY DO OYSTERS HAVE PEARLS  
WHY ARE DUCKS CALLED DUCKS  
WHY DO THEY CALL IT THE CLAP  
WHY ARE KYLE AND CARTMAN FRIENDS  
WHY IS THERE AN ARROW ON AANG'S HEAD  
WHY ARE TEXT MESSAGES BLUE  
WHY ARE THERE MUSTACHES ON CLOTHES  
WHY WUBA LUBBA DUB DUB MEANING  
WHY IS THERE A WHALE AND A POT FALLING  
WHY ARE THERE SO MANY BIRDS IN SWISS  
WHY IS THERE SO LITTLE RAIN IN WALLIS  
WHY IS WALLIS WEATHER FORECAST ALWAYS WRONG

WHY ARE THERE MALE AND FEMALE BIKES

WHY ARE THERE BRIDESMAIDS  
WHY DO DYING PEOPLE REACH UP  
HOW FAST IS LIGHTSPEED  
WHY ARE OLD KLINGONS DIFFERENT



WHY ARE THERE TINY SPIDERS IN MY HOUSE  
WHY DO SPIDERS COME INSIDE  
WHY ARE THERE HUGE SPIDERS IN MY HOUSE  
WHY ARE THERE LOTS OF SPIDERS IN MY HOUSE  
WHY ARE THERE SPIDERS IN MY ROOM  
WHY ARE THERE SO MANY SPIDERS IN MY ROOM  
WHY DO SPYDER BITES ITCH  
WHY IS DYING SO SCARY

WHY IS THERE NO GPS IN LAPTOPS  
WHY DO KNEES CLICK  
WHY AREN'T THERE 5 GRADES

WHY IS THERE CAFFEINE IN MY SHAMPOO  
WHY HAVE DINOSAURS NO FUR

WHY DO IGUANAS DIE

WHY AREN'T ECONOMISTS RICH  
WHY DO AMERICANS CALL IT SOCCER  
WHY ARE MY EARS RINGING  
WHY IS 42 THE ANSWER TO EVERYTHING  
WHY CAN'T NOBODY ELSE LIFT THORS HAMMER  
WHY IS MARVIN ALWAYS SO SAD

WHY ARE THERE ANTS IN MY LAPTOP

WHY IS EARTH TILTED  
WHY IS SPACE BLACK  
WHY IS OUTER SPACE SO COLD  
WHY ARE THERE PYRAMIDS ON THE MOON  
WHY IS NASA SHUTTING DOWN

WHY ARE THERE FEMALE N  
WHY ARE THERE TINY SPIDERS IN MY HOUSE  
WHY DO SPIDERS COME INSIDE  
WHY ARE THERE HUGE SPIDERS IN MY HOUSE  
WHY ARE THERE LOTS OF SPIDERS IN MY HOUSE  
WHY ARE THERE SPIDERS IN MY ROOM  
WHY ARE THERE SO MANY SPIDERS IN MY ROOM  
WHY DO SPYDER BITES ITCH  
WHY IS DYING SO SCARY  
WHY IS THERE NO GPS IN LAPTOPS  
WHY DO KNEES CLICK  
WHY AREN'T THERE 5 GRADES

WHY ARE SWISS AFRAID OF DRAGONS

# QUESTIONS

CAN BE ASKED BY ANYONE ANYTIME



WHY IS THERE AN OWL IN MY BACKYARD  
WHY IS THERE AN OWL OUTSIDE MY WINDOW  
WHY IS THERE AN OWL ON THE DOLLAR BILL  
WHY DO OWLS ATTACK PEOPLE  
WHY ARE FPGA'S EVERYWHERE  
WHY ARE THERE HELICOPTERS CIRCLING MY HOUSE  
WHY ARE THERE GODS  
WHY ARE THERE TWO SPOCKS

WHAT IS <https://xkcd.com/1256/>  
WHY DO THEY SAY T-MINUS  
WHY ARE THERE OBELISKS  
WHY ARE WRESTLERS ALWAYS WET  
WHY ARE OCEANS BECOMMING MORE ACIDIC

WHY IS THERE A LINE THROUGH HTTPS  
WHY IS THERE A RED LINE THROUGH HTTPS ON TWITTER  
WHY IS HTTPS IMPORTANT



WHY ARE THERE SO MANY CROWS IN ROCHESTER  
WHY IS TO BE OR NOT TO BE FUNNY  
WHY DO CHILDREN GET CANCER  
WHY IS POSEIDON ANGRY WITH ODYSSEUS  
WHY IS THERE ICE IN SPACE


WHY ARE DOGS AFRAID OF FIRE  
WHY IS THERE NO KING IN EN  
WHY ARE MY BOOBS ITCHY  
WHY ARE CIGARETTES LEGAL  
WHY ARE THERE DUCKS IN MY POOL  
WHY IS JESUS WHITE  
WHY IS THERE LIQUID IN MY EAR  
WHY DO Q TIPS FEEL GOOD  
WHY DO PEOPLE DIE  
WHY AREN'T THERE GUNS IN HARRY POTTER



**Hes·so**  **VALAIS  
WALLIS**



Haute Ecole d'Ingénierie  
Hochschule für Ingenieurwissenschaften



Silvan Zahno [silvan.zahno@hevs.ch](mailto:silvan.zahno@hevs.ch)